

Probability Distributions on Structured Objects

September 17, 2013

Reminder


- HW1 is due at 11:59pm tonight
- There was some ambiguity in this assignment
- The TAs gave a lot of help, but in general, learning to work from incomplete specs is important

Probability Outline

- Why probability?
- Probability review
- Multinomials vs. exponential parameterization
- Locally vs. globally normalized models & partition functions
- Examples

Why Probability?

- Probability formalizes
 - The concept of **models**
 - The concept of **data**
 - The concept of **learning**
 - The concept of **prediction** (inference)



Probability is expectation founded upon partial knowledge.

Why Probability?


- What might we have partial knowledge about?
 - The state of the world (test data)
 - The reliability of our training data
 - The correctness of our model
 - The values of our parameters

$$p(x \mid \text{partial knowledge})$$

What is a Probability?

- **Limiting (relative) frequency of events**
 - in repeated (identical) experiments
- **Degree of belief**
 - Subjective conception
 - 40% chance of rain tomorrow in Pittsburgh
- Viewpoint affects
 - interpretation
 - **not** rules of probability calculus themselves

Discrete Distributions



Discrete distribution: Ω is *finite* or *countable*, but no bigger

Discrete Distributions

$$\forall x \in \Omega, f(x) \in [0, 1]$$

$$\sum_{x \in \Omega} f(x) = 1$$

Probability mass function


An **event** is a subset (maybe one element) of the sample space, $E \subseteq \Omega$

$$P(E) = \sum_{x \in E} f(x)$$

Random Variables

A **random variable** is a function from a random event from a set of possible outcomes (Ω) and a probability distribution (ρ), a function from outcomes to probabilities.

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$\rho_X(x) = \begin{cases} \frac{1}{6} & \text{if } x = 1, 2, 3, 4, 5, 6 \\ 0 & \text{otherwise} \end{cases}$$


Random Variables

A **random variable** is a function from a random event from a set of possible outcomes (Ω) and a probability distribution (ρ), a function from outcomes to probabilities.

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

$$Y(\omega) = \begin{cases} 0 & \text{if } \omega \in \{2, 4, 6\} \\ 1 & \text{otherwise} \end{cases}$$

$$\rho_Y(y) = \begin{cases} \frac{1}{2} & \text{if } y = 0, 1 \\ 0 & \text{otherwise} \end{cases}$$



Sampling Notation

$$x = 4 \times z + 1.7$$



Variable

Expression

Sampling Notation

$$x = 4 \times z + 1.7$$

$$y \sim \text{Distribution}(\theta)$$

Random variable **Distribution** *Parameter*



Sampling Notation

$$x = 4 \times z + 1.7$$

$$y \sim \text{Distribution}(\theta)$$

$$y' = y \times x$$



Random variable

Joint Probability

- Probability over multiple event types
- Tool for reasoning about dependent (correlated) events

A **joint probability distribution** is a probability distribution over r.v.'s with the following form:

$$Z = \begin{bmatrix} X(\omega) \\ Y(\omega) \end{bmatrix}$$

$$\sum_{x \in \mathcal{X}, y \in \mathcal{Y}} \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) = 1 \quad \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) \geq 0 \quad \forall x \in \mathcal{X}, y \in \mathcal{Y}$$

Joint Probability

- Probability over multiple event types
- Tool for reasoning about dependent (correlated) events

A **joint probability distribution** is a probability distribution over r.v.'s with the following form:

$$Z = \begin{bmatrix} X(\omega) \\ Y(\omega) \end{bmatrix}$$

Words Tags

$$\sum_{x \in \mathcal{X}, y \in \mathcal{Y}} \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) = 1 \quad \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) \geq 0 \quad \forall x \in \mathcal{X}, y \in \mathcal{Y}$$

Joint Probability

- Probability over multiple event types
- Tool for reasoning about dependent (correlated) events

A **joint probability distribution** is a probability distribution over r.v.'s with the following form:

$$Z = \begin{bmatrix} X(\omega) \\ Y(\omega) \end{bmatrix}$$

Words
Trees

$$\sum_{x \in \mathcal{X}, y \in \mathcal{Y}} \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) = 1 \quad \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) \geq 0 \quad \forall x \in \mathcal{X}, y \in \mathcal{Y}$$

Joint Probability

- Probability over multiple event types
- Tool for reasoning about dependent (correlated) events

A **joint probability distribution** is a probability distribution over r.v.'s with the following form:


$$Z = \begin{bmatrix} X(\omega) \\ Y(\omega) \end{bmatrix}$$

DNA sequence Proteins


$$\sum_{x \in \mathcal{X}, y \in \mathcal{Y}} \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) = 1 \quad \rho_Z \left(\begin{bmatrix} x \\ y \end{bmatrix} \right) \geq 0 \quad \forall x \in \mathcal{X}, y \in \mathcal{Y}$$

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

$$X(\omega) = \omega$$





$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$


$$\begin{aligned}\Omega = & \{(1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6), \\& (2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6), \\& (3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6), \\& (4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6), \\& (5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6), \\& (6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6), \}\end{aligned}$$



$$X(\omega) = \omega_1 \quad Y(\omega) = \omega_2$$



$$\rho_{X,Y}(x,y) = \begin{cases} \frac{1}{36} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$



$$X(\omega) = \omega$$

$$\begin{aligned}\Omega = & \{(1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6), \\& (2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6), \\& (3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6), \\& (4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6), \\& (5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6), \\& (6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6)\}\end{aligned}$$



$$X(\omega) = \omega_1 \quad Y(\omega) = \omega_2$$

$$\rho_{X,Y}(x,y) = \begin{cases} \frac{x+y}{252} & \text{if } (x,y) \in \Omega \\ 0 & \text{otherwise} \end{cases}$$


Marginal Probability

$$p(X = x, Y = y) = \rho_{X,Y}(x, y)$$

$$p(X = x) = \sum_{y' \in \mathcal{Y}} p(X = x, Y = y')$$


$$p(Y = y) = \sum_{x' \in \mathcal{X}} p(X = x', Y = y)$$

$$\Omega = \{(1, 1), (1, 2), (1, 3), (1, 4), (1, 5), (1, 6), \\(2, 1), (2, 2), (2, 3), (2, 4), (2, 5), (2, 6), \\(3, 1), (3, 2), (3, 3), (3, 4), (3, 5), (3, 6), \\(4, 1), (4, 2), (4, 3), (4, 4), (4, 5), (4, 6), \\(5, 1), (5, 2), (5, 3), (5, 4), (5, 5), (5, 6), \\(6, 1), (6, 2), (6, 3), (6, 4), (6, 5), (6, 6), \}$$



$$p(X = 4) = \sum_{y' \in [1, 6]} p(X = 4, Y = y')$$


$$p(Y = 3) = \sum_{x' \in [1, 6]} p(X = x', Y = 3)$$


Marginal Probability




Marginal Probability



Marginal Probability



Marginal Probability



Marginal Probabilities

- In a joint model of word and tag sequences $p(w,t)$
 - The probability of a word sequence $p(w)$
 - The probability of a tag sequence $p(t)$
 - The probability of a word sequence with the word “cat” somewhere in it
 - The probability of a tag sequence containing three verbs in a row

Conditional Probability

The **conditional probability** is defined as follows:

$$p(X = x \mid Y = y) = \frac{p(X = x, Y = y)}{p(Y = y)} = \frac{\text{joint probability}}{\text{marginal}}$$

This assumes $p(Y = y) \neq 0$

We can construct joint probability distributions out of conditional distributions:


$$p(x \mid y)p(y) = p(x, y) = p(y \mid x)p(x)$$

Conditional Probability Distributions


The **conditional probability distribution** of a variable X given a variable Y has the following properties:

$$\forall y \in Y, \sum_{x \in X} p(X = x \mid Y = y) = 1$$

Conditional Probability



Conditional Probability




Conditional Probabilities

- In a joint model of word and tag sequences $p(w,t)$
 - The probability of a tag sequence given a word sequence $p(t | w)$
 - The probability of a word sequence given a tag sequence $p(w | t)$

Joint and Marginal Probabilities

- In a joint model of word and tag sequences $p(w, t)$
 - The probability that the 3rd tag is VERB, given $w = \text{"Time flies like an arrow"}$
 $p(t_3 = \text{VERB} | w = \text{Time flies like an arrow})$
 - The probability that the 3rd word is like, given $w = \text{"Time flies _____ an arrow"}$, $t_3 = \text{VERB}$
 $p(t_3 = \text{like} | w = \text{Time flies _____ an arrow}, t_3 = \text{VERB})$



Chain Rule

$$\begin{aligned} p(a, b, c, d, \dots) = & p(a) \times \\ & p(b \mid a) \times \\ & p(c \mid a, b) \times \\ & p(d \mid a, b, c) \times \\ & \vdots \end{aligned}$$

Bayes Rule

$$p(x | y) = \frac{p(y | x)p(x)}{p(y)} \quad \left(= \frac{p(y | x)p(x)}{\sum_{x'} p(y | x')p(x')} \right)$$

Diagram illustrating the components of Bayes Rule:

- Posterior**: Points to the term $p(x | y)$.
- Likelihood**: Points to the term $p(y | x)$.
- Prior**: Points to the term $p(x)$.
- Evidence**: Points to the term $p(y)$.

Independence

Two r.v.'s are **independent** iff

$$p(X = x, Y = y) = p(X = x) \times p(Y = y)$$

Equivalently (prove with def. of cond. prob.)

$$p(X = x \mid Y = y) = p(X = x)$$

Alternatively,

$$p(Y = y \mid X = x) = p(Y = y)$$

Conditional Independence

Two equivalent statements of conditional independence:

$$p(a, c | b) = p(a | b)p(c | b)$$

and:

$$p(a | b, c) = p(a | b)$$

“If I know B, then C doesn’t tell me about A”

$$p(a | b, c) = p(a | b)$$

$$p(a, b, c) = p(a | b, c)p(b, c)$$

$$= p(a | b, \cancel{c})p(b | c)p(c)$$

Conditional Independence

Two equivalent statements of conditional independence:

$$p(a, c | b) = p(a | b)p(c | b)$$

and:

$$p(a | b, c) = p(a | b)$$

“If I know B, then C doesn’t tell me about A”

$$p(a | b, c) = p(a | b)$$

$$\begin{aligned} p(a, b, c) &= p(a | b, c)p(b, c) \\ &= p(a | b, \cancel{c})p(b | c)p(c) \\ &= p(a | b)p(b | c)p(c) \end{aligned}$$

Conditional Independence

- Useful thing to assume when designing models
 - Limit the variables that influence distributions
 - Classical example: Markov assumption
- Questions
 - Does conditional independence imply marginal independence?
 - Does marginal independence imply conditional independence?

Expected Values

$$\mathbb{E}_{p(X=x)} [f(x)] \doteq \sum_{x \in \mathcal{X}} p(X = x) \times f(x)$$


Some special expectations:

$$p(X = y) = \mathbb{E}_{p(X=x)} [\mathbb{I}_{x=y}]$$

$$H(X) = \mathbb{E}_{p(X=x)} [-\log_2 x]$$

Categorical (Multinomial) Distributions

- Generalized model of a di to k dimensions
- Option 1: Parameters lie on the **k -simplex**



Log-linear Parameterization

Weight vector

Feature vector function

$$p(x) = \frac{\exp \mathbf{w}^\top \mathbf{f}(x)}{Z}$$

$$\text{where } Z = \sum_{x' \in \mathcal{X}} \exp \mathbf{w}^\top \mathbf{f}(x')$$

Assumption: Z converges

Categorical (Multinomial) Distributions


- “Naïve” parameterization
 - k outcomes, $k(-1)$ independent parameters
 - Model as tables of (conditional) probabilities
 - MLE estimation (given fully observed data) is easy
- Log-linear parameterization
 - k outcomes, n , possibly overlapping parameters
 - Share statistical strength across “related” events
 - How are elements related? Depends how you define f


Locally Normalized Models

- Structure as the result of a **discrete time branching process**
 - Start in a known initial state, carry out stochastic steps (parameterized using multinomials) until some termination condition is met
 - Steps are (conditionally) independent of one another: probabilities multiply
 - *Total probability is the probability of the steps*
- Usually for joint (generative) models
 - not always though (see Appendix D.2)


S

1.0



$$1.0 \times p(\text{NP VP} \mid \text{S})$$




$$1.0 \times p(\text{NP VP} \mid \text{S}) \\ \times p(\text{JJ NN} \mid \text{NP})$$




$$\begin{aligned} & 1.0 \times p(\text{NP VP} \mid \text{S}) \\ & \times p(\text{JJ NN} \mid \text{NP}) \\ & \times p(\text{V} \mid \text{VP}) \end{aligned}$$




$1.0 \times p(\text{NP VP} \mid \text{S})$
 $\times p(\text{JJ NN} \mid \text{NP})$
 $\times p(\text{V} \mid \text{VP})$
 $\times p(\text{angry} \mid \text{JJ})$



$1.0 \times p(\text{NP VP} \mid \text{S})$
 $\times p(\text{JJ NN} \mid \text{NP})$
 $\times p(\text{V} \mid \text{VP})$
 $\times p(\text{angry} \mid \text{JJ})$
 $\times p(\text{dogs} \mid \text{NN})$



$$p(\tau, \mathbf{x}) = \prod_{r \in \mathcal{G}} p(r \mid \mathcal{G})^{f(r \in \tau)}$$



$1.0 \times p(\text{NP VP} \mid \text{S})$
 $\times p(\text{JJ NN} \mid \text{NP})$
 $\times p(\text{V} \mid \text{VP})$
 $\times p(\text{angry} \mid \text{JJ})$
 $\times p(\text{dogs} \mid \text{NN})$
 $\times p(\text{bark} \mid \text{V})$

Here's an alternative way of building a tree and string:

S


1.0

Here's an alternative way of building a tree and string:




$$1.0 \times p(2 \text{ kids} \mid S)$$

Here's an alternative way of building a tree and string:




$$1.0 \times p(\text{2 kids} \mid S) \\ \times p(\text{NP} \mid S, n=1, \text{total}=2)$$

Here's an alternative way of building a tree and string:




$$\begin{aligned} & 1.0 \times p(2 \text{ kids} \mid S) \\ & \times p(\text{NP} \mid S, n=1, \text{total}=2) \\ & \times p(\text{VP} \mid S, n=2, \text{total}=2) \end{aligned}$$

Here's an alternative way of building a tree and string:




1.0 x $p(2 \text{ kids} \mid S)$
x $p(\text{NP} \mid S, n=1, \text{total}=2)$
x $p(\text{VP} \mid S, n=2, \text{total}=2)$
x $p(1 \text{ kid} \mid \text{VP})$

Here's an alternative way of building a tree and string:




Here's an alternative way of building a tree and string:



$$\begin{aligned} & 1.0 \times p(2 \text{ kids} \mid S) \\ & \times p(\text{NP} \mid S, n=1, \text{total}=2) \\ & \times p(\text{VP} \mid S, n=2, \text{total}=2) \\ & \times p(1 \text{ kid} \mid \text{VP}) \end{aligned}$$

Here's an alternative way of building a tree and string:



Choosing a Model

- Independence is a property of distributions
 - Look at distributions in the wild, figure out what independence assumptions hold
- Dependence makes modeling more expensive
 - How big does your CKY chart have to be if you have “grandparent” annotation?

Parameterization

- For each step in the branching process
 - We have a multinomial distribution
 - We can use independent parameters (on simplex)
 - We can use log-linear models
 - “Locally normalized model” (cf. Appendix D.2)
 - Z is “local” to the decision being made

Globally Normalized Models

- Extension of the exponential parameterization to structured output spaces

$$p(\mathbf{x}) = \frac{\exp \mathbf{w}^\top \mathbf{F}(\mathbf{x})}{Z}$$

$$\text{where } Z = \sum_{\mathbf{x}' \in \mathcal{X}} \exp \mathbf{w}^\top \mathbf{F}(\mathbf{x}')$$

Conditional Random Fields

$$p(\mathbf{y} \mid \mathbf{x}) = \frac{\exp \mathbf{w}^\top \mathbf{F}(\mathbf{x})}{Z(\mathbf{x})}$$

$$Z(\mathbf{x}) = \sum_{\mathbf{y}' \in \mathcal{Y}_x} \exp \mathbf{w}^\top \mathbf{F}(\mathbf{x})$$

Conditional Random Fields


$$p(\mathbf{y} \mid \mathbf{x}) = \frac{\exp \mathbf{w}^\top \mathbf{F}(\mathbf{x}, \mathbf{y})}{Z(\mathbf{x})}$$

$$Z(\mathbf{x}) = \sum_{\mathbf{y}' \in \mathcal{Y}_x} \exp \mathbf{w}^\top \mathbf{F}(\mathbf{x}, \mathbf{y}')$$

Decoding is nice:

$$\begin{aligned}\mathbf{y}^* &= \arg \max_{\mathbf{y} \in \mathcal{Y}_x} \frac{\exp \mathbf{w}^\top \mathbf{F}(\mathbf{x}, \mathbf{y})}{Z(\mathbf{x})} \\ &= \arg \max_{\mathbf{y} \in \mathcal{Y}_x} \exp \mathbf{w}^\top \mathbf{F}(\mathbf{x}, \mathbf{y}) \\ &= \arg \max_{\mathbf{y} \in \mathcal{Y}_x} \mathbf{w}^\top \mathbf{F}(\mathbf{x}, \mathbf{y})\end{aligned}$$

Conditional Random Fields



$$F(x, y) = \sum_{C \in G} f(C)$$

Comparison of Feature-Based Models

- Locally Normalized Models
 - Good joint models
 - Easy to training
 - Downside: decoding can be expensive
- Globally Normalized Models
 - Very popular conditional models (CRFs)
 - Challenge: computing Z / training
 - Advantage: decoding can be cheap